

Bioinformatics Algorithms

Week 1

Stephan Peischl

University of Bern

This course

12 lectures, 90 minutes each.

HS 2026: **14 September – 18 December**.

Two calendar weeks are buffer. No 13th lecture.

Language: **English** (exam in English or German). Slides on **ILIAS**. Code is **Python**.

Lab: one Jupyter notebook each week.

Open in **VS Code**: Select Kernel, then **Run All**.

Questions: stephan.peischl@unibe.ch ·
D401, Baltzerstrasse 6

Prerequisites: basic math and some programming. Biology helps as context.

Exam

Written, 2 hours, closed book · date TBA

Sign up on **KSL** at least 14 days before.

You **outline an idea**. No code.

Each lecture ends with three check questions.

Q1–Q2 test today's idea.

Q3 is a situation we did not practise.

If you can follow that week's lab notebook *and* answer Q3, you are in exam shape.

Models can write the code. That is not the course.

A language model can emit a Fibonacci function, or a Hanoi solver, in seconds. That is not the same as **having an algorithm you can check**.


Is it **correct** for every n ?

What does it **cost** as n grows — time, memory, and how long the output is?

Is there a **better** method?

Generated code is a guess until you can answer those three questions. That is what this course trains.

A test of “reasoning”: Towers of Hanoi



Rules: move one disk at a time; never put a larger disk on a smaller one.

The shortest solution for n disks is not a long checklist. It is **the same puzzle, smaller**, twice, with one move in the middle.

Hanoi $n = 3$, live


Try to solve the puzzle for 3 disks on paper. Can you find a general way to solve this puzzle?

Hint: Watch for the pattern: a pile of two goes to the spare peg, the largest disk moves once, the pile of two comes back.

Three jobs

To get n disks to the goal: move the top $n - 1$ **out of the way**, move the largest disk **once**, move the $n - 1$ **on top**. Each “move $n - 1$ ” is the same puzzle, one disk smaller.

To move n disks: same puzzle on $n-1$, one move, same puzzle on $n-1$.



$T(n) = T(n-1) + 1 + T(n-1)$. One disk: $T(1) = 1$. Then 1, 3, 7, 15, ... = $2^n - 1$.

For $n = 3$: 3 + 1 + 3 moves. The same three jobs work for any n .

The recursive algorithm

Those three jobs, typed in. Same code as in the lab notebook.

```
1 def hanoi(n, src, dst, aux):
2     if n == 0:
3         return
4     hanoi(n - 1, src, aux, dst)    # smaller tower to aux
5     print(src, "->", dst)         # largest disk once
6     hanoi(n - 1, aux, dst, src)    # smaller tower to goal
```

`hanoi(3, "start", "goal", "aux")` prints 7 lines. The function body is three steps; the **printout** is what grows.

The call tree for $n = 3$

Each call does two smaller calls and one printed move. Unfold that for $n = 3$:

The algorithm text stays short. The tree of calls — and the list of moves — doubles with every extra disk: $T(n) = 2^n - 1$.

An iterative version

You can solve Hanoi **without recursion**. Same $2^n - 1$ moves; a loop instead of a call stack.

Rule (odd n). Repeat until done:

1. Move the **smallest** disk one step in the cycle `start → goal → aux → start`.
2. Make the **only other legal move** (the move that does not touch the smallest disk).

For even n , reverse the cycle: `start → aux → goal → start`.

```
1 # Pseudocode – same output length as the recursive version
2 for step in range(1, 2**n):          # 2^n – 1 moves
3     if step % 2 == 1:
4         move_smallest_along_cycle()
5     else:
6         make_only_other_legal_move()
```

Recursive or iterative: the **idea** can be short. Emitting every move still costs $2^n - 1$ lines.

The list of moves is the expensive part

The algorithm is **tiny**: three jobs (or a short loop). The *list of moves* is $2^n - 1$ lines.

Printing every move is a resource, like time and memory.

Disks n	Moves $2^n - 1$
3	7
8	255
10	1,023
20	$\approx 10^6$

Describing the algorithm is not the same job as emitting every move. The list has its own cost.

Every legal board is a vertex



Each disk can sit on any of three pegs, with the size rule. That gives 3^n **legal boards**. Draw a vertex for each board and an edge for each legal one-disk move: for small n the picture is a **Sierpiński triangle**.

A shortest solution is a path of length $2^n - 1$ from “all on start” to “all on goal.” The three-line recurrence *compresses* that path; you do not need to search the graph.

Week 7 draws Hanoi as a graph properly (nodes, edges, walks). Today: the state space is huge, fractal-shaped, and the recurrence already knows a shortest path.


Can LLMs reason?

Shojaee et al. (2025) asked frontier models for Hanoi **move lists** as n grew. The same models could often **state** the three jobs — and still fail to carry them out.



Accuracy is already gone by $n \approx 8$, where $2^8 - 1 = 255$ lines still fit. That is not “they ran out of words.”

Collapse is not special to Hanoi



Thinking models last a few extra steps. Then **everyone** is at zero — Hanoi, checkers, blocks, river crossing.


Shojaee et al., Fig. 5. Shaded bands: low / medium / high complexity. [arXiv:2506.06941](https://arxiv.org/abs/2506.06941). *Huge n* also hits output caps (Opus & Lawsen) — a different story.

Even given the algorithm

Hanoi has an **exact** solution: the same three jobs, again.

They put that algorithm **in the prompt**.
The model only had to execute the steps.
Collapse still happens at the same n .

A language model does not *run* the three jobs. It *re-predicts* a plausible next board, move after move. One illegal move poisons the rest.



Name the problem first, then pick a method that actually carries out the steps.

What we will mean by algorithm

An **algorithm** is a finite list of well-defined instructions that starts from a given input, passes through a finite sequence of states, and **terminates** with an output.

Hanoi's three jobs are an algorithm. A random mix of moves that might never finish is not.

Not a recipe

1. Take ingredients.
2. Mix them in a random way.
3. Bake at a random temperature.
4. If the result is a cupcake: stop.
5. Else: go to 1.

Steps are not well-defined; it need not terminate.

Three questions, every week: is it **correct**? what does it **cost** as a function of n (time, memory,

Input and output

Every algorithm in this course will be stated as:

INPUT

What you are given. Size measured by n (and maybe m).



OUTPUT

What you must produce. A number, a path, a labelling, a tree.

If you cannot name input and output, you do not yet have a problem — only a vibe.

Three questions, every week

We study **how algorithms work**: the idea, a worked example, what it costs.

Is it **correct**?

What does it **cost** as a function of n ? Time, memory, *and how long the answer is*.

Can we do **better**?

Every meeting has a **lab notebook**: the lecture example, already filled in. Run it, change an input, watch the printout. We close with history, tools that use the idea, and a short reading list.

Twelve weeks, in plain language

Biology gives us the data (strings, graphs, trees). This course is about **how to organise that data** so a method becomes feasible — not about memorising software menus.

Weeks	What we do
1–2	Learn to judge a method (correct? $T(n)$? better?). Meet problems where listing every answer fails.
3–6	Find similarity between sequences (alignment as a filled table). Search smarter by pruning hopeless branches (motifs, branch-and-bound). Ask when a greedy step is good enough.
7–9	Combine pieces into one genome (assembly: change the graph). Search a long string without scanning it all (suffix trees, BWT / an index).
10–12	Reconstruct hidden structure : trees (phylogenetics), groups or axes (clustering / PCA), a classifying line (SVM).

Change how the data are organised — a table, a tree of partial answers, a graph, an index —

and the algorithm changes with them.

Week-by-week titles live on ILIAS with the slides. We study algorithms, not software manuals.

Why size forces a choice

A genome is a long string. What we measure are short pieces.

ATCGTTGACCGGAA | $\sim 3 \times 10^9$ letters (human)

TTGACCGG ACCGGAAT GAATCGTT

Questions this course turns into algorithms

How similar are two strings? How do we assemble fragments? How do we search 3×10^9 letters?

A procedure that is exponential in n is already dead at genome scale.

A human genome is about 3×10^9 letters. Sequencing gives **short fragments**. Alignment, assembly, and search are procedures on strings of this size — that is weeks 4, 7, and 8–9, not a different subject.

All-vs-all character comparisons on that n would be $n^2 \approx 9 \times 10^{18}$ operations. At a nanosecond each, that is not “a slow laptop” — it is **centuries**. Later we will change how the letters are organised (an index, a graph), not buy a faster CPU.

An exponential — or even quadratic — procedure is not “a bit slow”. At genome scale it does not finish. Change the algorithm, or change how the data are organised.

Today's toy example

1. State input, output, and the three questions.
2. Explain why naive recursive Fibonacci is exponential, and how storing results makes it linear.
3. Read Big-O off the loops in selection sort and bubble sort.

Toy. $F_0 = 0, F_1 = 1, F_n = F_{n-1} + F_{n-2}$. Compute F_5 on paper, by unfolding the definition. (You should get 5.)

Same idea as Hanoi: name the output, then pick a method that gives you the answer. At the end you will run two programs in Python and time them.

Fibonacci as a specification

Fibonacci rabbits (Leonardo of Pisa): mature in one month, then one new pair each month



$$F_n = F_{n-1} + F_{n-2}$$

Input: n (months). Output: F_n (pairs). Today's toy: $n = 5$ gives 5.

INPUT

An integer $n \geq 0$ (months, or just an index).



OUTPUT

The integer F_n .

The obvious program

The recurrence, typed in:

```
1 def fib(n):  
2     if n < 2:  
3         return n  
4     return fib(n - 2) + fib(n - 1)  
5  
6 fib(5)      # 5  - the toy  
7 fib(10)     # 55
```

It matches the definition. For large n it becomes unusable.

Unfold F_5 on the board

Do **not** look at the next picture yet.

Write the call tree: F_5 calls F_4 and F_3 . Keep going until F_1 and F_0 .

Vote: how many times does F_3 appear?

Three questions, on this program

Is it **correct**?

How much **time** does it take, as a function of n ?

Can we **do better**?

Is it correct?

Yes. Each call returns $F_n = F_{n-1} + F_{n-2}$, with the right base cases.

How much **time** does it take, as a function of n ?

Can we **do better**?

How much time?

Let $T(n)$ be the number of basic operations to compute F_n with `fib`.

$$T(n) = T(n - 1) + T(n - 2) + \Theta(1)$$


$T(n)$ grows at least as fast as F_n itself.

$$T(200) \geq F_{200} \geq 2^{138} \geq 10^{41}$$

Correct, and useless for large n . Same lesson as the genome: what matters is how the cost **grows**, not how fast the laptop is.

Why it explodes

The recursive program does not remember earlier answers. When it needs F_3 again, it rebuilds the whole subtree.



F_3 twice, F_2 three times, F_1 three times — each time computed from scratch.

If we remembered each answer, we would not redraw those subtrees.

In this tree for F_5 : F_3 appears **twice**, F_2 **three** times, and F_1 **three** times. Same integers, computed again from scratch each time. That wasted work piles up, so the number of calls grows like F_n itself — exponentially.

Can we do better?

The explosion comes from **recomputation**: the same F_k is asked for again and again. So keep each value once you have it.

On the toy: the call tree for F_5 makes **15** calls — you already saw F_3 appear twice. If we store answers as we go, we only need **five additions** and never recompute. Both methods still return the integer 5.

Compute $F_1, F_2, \dots, F_n$ once each, in order. Each step looks up the two previous values.

Store previous values


How this works. Keep only the last two Fibonacci numbers, `a` and `b`. Each loop step advances one index: the new pair is `(b, a+b)`. After $n - 1$ steps, `b` is F_n . No tree, no recomputation — just a walk from 0 to n .

```
1 def fib1(n):
2     if n < 2:
3         return n
4     a, b = 0, 1
5     for _ in range(n - 1):
6         a, b = b, a + b
7     return b
8
9 fib1(5)    # same toy, linear work
```

Two costs

Algorithm	Idea	Cost
<code>fib</code>	Recurse on the definition	$T(n) \approx 2^n$
<code>fib1</code>	Remember previous values	$T(n) \approx n$

Same input/output, two algorithms. The difference is an **idea** (reuse), not a faster computer.



Measured on a linear axis. Recursive `fib` is already slow by $n \approx 30$. The stored-value version sits on the axis.

Performance is not complexity

Performance

How long this run took, on this machine.

Complexity

How the resource need **scales** as n grows.


A faster laptop does not change $O(2^n)$ into $O(n)$. That was the genome slide, in one line.

Big-O, informally

$O(f(n))$: for large n , the running time is bounded by a constant times $f(n)$.

$O(n^2)$ covers both $0.0005 n^2$ and $10^7 n^2 + 3n + 2$.

Only the **leading term** matters. Constant factors do not.



Formally

$f(n)$ is $O(g(n))$ if there exist $c > 0$ and n_0 such that $f(n) \leq c g(n)$ for all $n > n_0$.

Claim: $6n^4 - 2n^3 + 5$ is $O(n^4)$. For $n \geq 1$:

$$f(n) \leq 6n^4 + 2n^3 + 5 \leq 13n^4.$$

Rules we will actually use

- Drop constants: $O(k f(n)) = O(f(n))$
- Drop slower terms: $O(n^2 + n^3) = O(n^3)$

Notation	Name
$O(1)$	constant
$O(\log n)$	logarithmic
$O(n)$	linear
$O(n^2)$	quadratic
$O(n^k)$	polynomial
$O(k^n)$	exponential

Reading complexity off sorting

Same three questions, new output: a **sorted list**. We will not write a full sorter from scratch. We will read Big-O off the **loops that sorting uses**.

A **comparison** (`a[i] < a[j]`) or a **swap** of two entries is $O(1)$: one step, independent of n .
Nested loops are where n^2 comes from.

To add costs: line up the pieces and keep the slowest. If a branch can exit early, quote the **worst** case — the input that makes the program do the most work.

One pass of selection sort

Ninety seconds. Do not shout. Write Big-O on paper.

What this does. Walk the list once. Keep an index `m` of the smallest value seen so far. At the end, `m` points at the minimum of all n items. That scan is the inner step of selection sort. Each comparison is $O(1)$.

```
1 m = 0
2 for j in range(1, n):
3     if a[j] < a[m]:    # 0(1)
4         m = j
```

The loop runs about n times. Cost $O(n)$. One pass. Selection sort uses this to find the next item to place.

Selection sort: nested scans

How the algorithm works. Grow a sorted prefix from the left. For each position i , scan the unused suffix $a[i], \dots, a[n - 1]$, find its minimum, and **swap** that minimum into position i . Then i moves one step right. After n such placements, the whole list is sorted.

```
1 for i in range(n):
2     m = i
3     for j in range(i, n):
4         if a[j] < a[m]:           # 0(1) comparison
5             m = j
6     a[i], a[m] = a[m], a[i]      # 0(1) swap
```

Outer loop: n positions. Inner loop: up to n comparisons each. About $n^2/2$ comparisons in total — that is $O(n^2)$. Drop the constant. Bubble sort and insertion sort use the same nested shape.

Bubble sort: worst case, not the lucky pass


How the algorithm works. Repeatedly walk neighbouring pairs. Whenever two neighbours are out of order, swap them. One full walk is a **pass**. Large values “bubble” toward the end. If a pass makes **no** swap, the list is already sorted and we can stop. We still quote the **worst** case for Big-O.

```
1 swapped = True
2 while swapped:                # at most n passes
3     swapped = False
4     for i in range(n - 1):
5         if a[i] > a[i + 1]:    # 0(1)
6             a[i], a[i + 1] = a[i + 1], a[i]
7             swapped = True
```

Already sorted: one pass, $O(n)$. Reverse sorted: about n passes, $O(n^2)$. Big-O is the worst case, not the average over lucky inputs.

Those loops, as a picture

Same input, several methods. Selection and bubble are the nested $O(n^2)$ scans you just read. Merge sort and quicksort **split** the list and combine the pieces — that is where $O(n \log n)$ comes from. Watch them move.



More starting orders: Toptal, Sorting Algorithms.

A few sorting methods

Method	Idea	Typical	Worst case
Insertion, selection, bubble	Scan for the next item, or swap neighbours	$O(n^2)$	$O(n^2)$
Merge sort, heap sort	Split, sort halves, merge — or use a heap	$O(n \log n)$	$O(n \log n)$
Quicksort	Pick a pivot; recurse on the two sides	$O(n \log n)$	$O(n^2)$

Same output, different ideas, different growth — the same lesson as the two Fibonacci programs.

Back to the two Fibonacci

The recursion tree has at most $1 + 2 + 4 + \dots + 2^{n-1}$ nodes.

$$T_{\text{fib}}(n) = O(2^n), \quad T_{\text{fib1}}(n) = O(n)$$

Same output as the toy F_5 . Storing values turns exponential work into linear work.

Recursion is not the explosion

One last comparison before the lab. $n! = 1 \cdot 2 \cdot \dots \cdot n$. Same output, two programs:

```
1 def fact(n):                # n * (n-1)!
2     if n <= 1:
3         return 1
4     return n * fact(n - 1)   # one call
5
6 def fact1(n):                # 1*2*...*n
7     p = 1
8     for k in range(1, n + 1):
9         p *= k
10    return p
```

`fact` calls itself **once**. The calls are a chain of length n , not a binary tree.

`fib` is exponential because each call splits in **two**. Recursion itself is not the cost.

90 seconds: factorial vs Fib

Pair up. How many recursive calls does `fact(5)` make? How many does `fib(5)` make?

`fact(5)` is a chain of length 5. `fib(5)` is a tree. The wrap will ask this again as Q3.

The number of recursive *branches* is what matters, not the word recursion. That is how you recognise when an approach will work.

Open `labs/01-fib.ipynb`

LAB · ~15 MINUTES · PYTHON

The toy is still $F_5 = 5$. Run All — both programs are already in the notebook.

1. Confirm `fib(5) == 5` and `fib1(5) == 5`.
2. Time `fib` vs `fib1` for $n = 5, 10, 15, 20, 25$. Do not start at $n = 40$.
3. Why is `fib(25)` slower if both return the same integer?
4. `fact` vs `fact1` for $n!$: why is the recursive one not exponential?
5. Optional: `hanoi(3)` prints 7 lines. The algorithm is three lines; $n = 20$ would print about a million. Output length is a resource.

The code already runs. The questions are yours to answer.

A bit of history

- **al-Khwarizmi** (9th c.): the word *algorithm*. A finite, transmissible procedure.
- **Fibonacci**, *Liber Abaci* (1202): the rabbit recurrence. The algorithm is not the story.
- **Bellman** (1950s) will name what `fib1` is doing — fill a table of overlapping subproblems, later called dynamic programming (week 3).
- **Knuth**, *The Art of Computer Programming*: analysis of recurrences as a subject.

Tools that use this idea

- Every compiler and interpreter: you ask the same three questions of generated code — including code from a language model.
- Genome-scale tools exist because someone changed the *idea*, not the laptop. **BLAST** (week 4) indexes short words, then extends a local match. **minimap2** and **BWA** (weeks 8–9) index the genome so a query does not scan all 3×10^9 letters. **IQ-TREE** (week 10) scores trees by pruning, not by unfolding an exponential recurrence.
- Python Tutor — watch `fib(5)` allocate a tree of calls.

Extensions

- **Memoization.** Keep the recursive definition of `fib`, but store every answer in a dictionary the first time you compute it. The next call with the same n returns the stored value. Still $O(n)$ time (each F_k once), but the recursion depth can grow with n , so the call stack can overflow for large n . `fib1` avoids that by using a loop.
- **Matrix exponentiation.** There is a 2×2 matrix whose powers produce Fibonacci numbers. Raising that matrix to the n -th power needs only $O(\log n)$ matrix multiplications. Same output F_n , a third algorithm — faster asymptotically than `fib1`, but we will not need it this course.
- Recurrences in biology: branching processes; the number of histories on a tree (week 10) is not Fibonacci, but it is the same *kind* of count. Storing those histories as shared edges is a change of layout, not a new recurrence.

What you must be able to say

Without notes. Q1–Q2 check today's idea. Q3 is a situation we did not practise.

Q1 Why is `fib(25)` slower than `fib1(25)` if both return the same integer?

Q2 What is $T(n)$ for each, in Big-O?

Q3 $n! = 1 \cdot 2 \cdots n$. Recurrence: $n \cdot (n - 1)!$, one call. Why is `fact` $O(n)$, like the loop, not $O(2^n)$ like `fib`?

Q3 is a question we did not practise. Replaying today's toy is not enough.

Takeaways

Name the input and the output. Then ask: is it correct? what is $T(n)$? can we do better? Count **output length** as a resource too.

Naive Fibonacci recomputes the same values many times. Storing values turns $O(2^n)$ into $O(n)$.

Hanoi has three exact steps. Printing the list of moves is a different resource. Asking a language model for that list is the wrong tool — a guess at every move, not those three steps. At genome scale, **how the work grows** is what counts.

Next week

Hard problems and greedy ideas. Same three questions; the object is a map of cities, not a recurrence.

Toy: four cities. How many tours? A greedy hop to the nearest city — is it shortest?

To read

- Shojaee et al. (2025). *The Illusion of Thinking*. [arXiv:2506.06941](#). Comment: [Opus & Lawsen](#).
- Dasgupta, Papadimitriou, Vazirani. *Algorithms* — free PDF, ch. 0–2. [Berkeley page](#).
- Cormen, Leiserson, Rivest, Stein. *Introduction to Algorithms* — Big-O as used in this course.
- Step through `fib` in [Python Tutor](#).
- Sorting animations: [Toptal](#).