

Hard problems and greedy ideas

Week 2

Stephan Peischl

University of Bern

Recap

Last week's three questions: is it **correct**? what is $T(n)$? can we do **better**?

Same integer F_n , two algorithms: naive recursion $O(2^n)$, store previous values $O(n)$. Growth, not the laptop.

We also met a human genome: about 3×10^9 letters. Quadratic comparison already does not finish. Alignment, assembly, and search are procedures on **that** n .

Today we look at **TSP**: a classic problem that is hard and often turns up in disguise — for example in genome assembly.

Labwork doesn't give you a chromosome

You want one sequence: a pathogen in an outbreak, a tumour, a new isolate, a reference to call variants against.

What you hold is a **bag of short reads** — typically a few hundred letters each, millions of them. Until those pieces are ordered, you cannot read a gene or compare two genomes.

A genome is a long string. What we measure are short pieces.

ATCGTTGACCGGAA | $\sim 3 \times 10^9$ letters (human)

TTGACCGG

ACCGGAAT

GAATCGTT

Questions this course turns into algorithms


How similar are two strings? How do we assemble fragments? How do we search 3×10^9 letters?

A procedure that is exponential in n is already dead at genome scale.

Last week this picture was “the string is huge.” Today the bar at the top is **unknown**. The coloured pieces are what was measured.

Shotgun sequencing

DNA is sheared. The machine reads the fragments, not the intact chromosome.




Shotgun sequencing. National Human Genome Research Institute, Wikimedia Commons, public domain. File: Shotgun sequencing lg.jpg.

The fragments **overlap**. Reconstruction means putting them in an order so neighbours agree on the letters they share.

What the data look like

An unknown genome; four reads that tile it. Neighbours share a suffix–prefix. A wrong order leaves a gap.

The genome is unknown. The machine gives overlapping pieces.



A good order: neighbours overlap. A bad order: they do not.

Input: a set of overlapping reads. Output: an order that rebuilds one string. That is the biological problem.

A tiny genome: **ACAAT**

Suppose the hidden string is **ACAAT**. The machine might return **ACA**, **CAA**, **AAT**.

Assembly teaser. Fragments of ACAAT. Visit every read once.



ACA then CAA then AAT spells ACAAT. Every read used once.

Other orders of the three reads do not reconstruct the string.

ACA then **CAA** share **CA**. **CAA** then **AAT** share **AA**. That path visits every read **once** and spells **ACAAT**.

Six orders, one genome

Three reads can be lined up in $3! = 6$ orders. Only `ACA → CAA → AAT` reconstructs `ACAAT`. The other five disagree on the overlap.

Order	Overlaps?
<code>ACA</code> , <code>CAA</code> , <code>AAT</code>	yes — spells <code>ACAAT</code>
<code>ACA</code> , <code>AAT</code> , <code>CAA</code>	no
<code>CAA</code> , <code>ACA</code> , <code>AAT</code>	no
... three more ...	no

For n reads there are $n!$ orders (before symmetries). Most are wrong. Choosing the order is the algorithm — not a laboratory step.

Greedy assembly of reads

A natural rule: start at a read, then always glue the unused neighbour with the **largest overlap**.

On the toy: from `ACA`, `CAA` shares two letters, `AAT` shares one. Greedy takes `CAA`, then `AAT`.
Lucky — that is the true string.

That rule is **nearest neighbor**. It will fail on a bigger map. Today we need numbers we can add, so we rename the objects.

Travelling salesman problem

Visiting every **read** once is the same puzzle as visiting every **city** once and returning home.

Assembly

Reads. Distance = how badly they overlap (or $1 / \text{overlap}$). A path through every read rebuilds the genome.

A map

Cities. Distance = kilometres (or the lab matrix). A **tour** visits every city and returns home.

INPUT

n cities and a symmetric matrix of pairwise distances.



OUTPUT

A cyclic order of the n cities (a tour) of minimum total length.

TSP adds “return home.” The explosion is the same: the number of orders is factorial. Week 7 changes how assemblers draw the graph; today the working object is four cities and a distance

The idea

Exhaustive search is a method, not a solution: listing every tour of n cities is $O(n!)$. A greedy step is locally best and may be globally wrong — unless you can prove it. For some problems, a good approximation is the right goal.

By the end of this meeting you should be able to:

1. Count TSP tours, list them on $n = 4$, and score them.
2. Build a nearest-neighbor tour from **two different starting cities**, then try to improve one tour by **swapping two edges** (2-opt).
3. Run a genetic algorithm on tours (with parameters you can change), and earliest-finish on intervals.

The same tour, many labels

TSP is rarely advertised as “cities.” Anything that must **visit every object once and return** (or almost) is this problem.

Biology

Order overlapping **reads** (today). Order DNA probes on a chip. A robot spotting samples on a microarray.

Elsewhere

Drill every hole on a circuit board. Deliver to every address. Point a telescope at every target. A milling tool that must visit every cut.

The cities change; the factorial count does not. If you can name the objects and the cost of going from one to the next, you have TSP — often in disguise.

Recognise the tour, then pick a method. Do not wait for someone to say “travelling salesman.”

How many tours?

For n cities there are

$$\frac{(n-1)!}{2}$$

distinct tours: direction does not matter, and the starting city is only a rotation of the same cycle.

Four cities

A tour is a cycle. Starting at B on

$A \rightarrow B \rightarrow C \rightarrow D \rightarrow A$


is still that tour. Walking it backwards,

$A \rightarrow D \rightarrow C \rightarrow B \rightarrow A$,

is the same tour. That is why we count $(n - 1)!/2$.

Five cities: $4!/2 = 12$. Ten: 181 440.

Twenty: about 6×10^{16} .




Why $(n - 1)!/2$?

Fix a start, say A . The other $n - 1$ cities can be ordered in $(n - 1)!$ ways. Walking a cycle backwards is the same tour, so divide by 2.

n	Distinct tours
3	$2!/2 = 1$
4	$3!/2 = 3$
5	$4!/2 = 12$
6	$5!/2 = 60$

Three cities. One tour. Direction is the same cycle.



$(3-1)!/2 = 1$. Four cities: three tours. That is the formula.

Draw $n = 3$: only one triangle. $(3 - 1)!/2 = 1$. Four cities: three tours. That is the formula, not a mystery.

Three tours, scored


The numbers in the lab matrix are pairwise costs, **not** distances on a geographic map. Add the four edges of each cycle.

Four cities: three tours. Rotation and direction are the same tour.

distances


	A	B	C	D
A	0	2	6	3
B	2	0	3	1
C	6	3	0	4
D	3	1	4	0

A-B-C-D-A




$$2+3+4+3 = 12$$

A-B-D-C-A



$$2+1+4+6 = 13$$

A-C-B-D-A



$$6+3+1+3 = 13$$

List them. Score them. The shortest is 12. That is exhaustive search on $n=4$.


Exhaustive search

Method: generate every tour, keep the shortest.

```
1 best = float("inf")
2 for tour in all_tours(cities):
3     length = tour_length(tour, dist)
4     if length < best:
5         best = length
```

Correct. Cost: $\Theta(n!)$ tours, each $O(n)$ to score.

`fib` was exponential. Exhaustive TSP is **factorial** — worse. Listing every order of n reads is the same count.



How far does exact search go?

n	Tours $(n - 1)!/2$	Roughly
5	12	instant
10	1.8×10^5	a second
12	2.0×10^7	a minute
15	4.4×10^{10}	not today
20	6×10^{16}	no

A better exact method (Held–Karp, next slide) improves this to $O(n^2 2^n)$ — still exponential. Twenty cities is already hopeless for exact search in class.

Exact TSP can reuse overlapping paths

Held–Karp (1962) does **not** list every tour. It fills a table.

What one cell means. Fix a start city A . A cell stores the length of the **best path** that

- started at A ,
- visited exactly the cities in a set S ,
- and **ended** at a particular city $i \in S$.

When you grow S , you reuse the answers already stored for smaller sets — the same diagnosis as Fibonacci.

That method — fill a table of overlapping subproblems — is **dynamic programming**. It is the topic of **next week**.

About $n \cdot 2^n$ cells, not $n!$ tours. For $n = 20$: roughly $20 \cdot 10^6$ cells vs 6×10^{16} tours — better, still not something we fill by hand.

Some problems stay expensive

Last week: two algorithms, two costs, **same problem**.

Some problems are expensive **no matter how clever the algorithm**. Listing all subsets of an n -element set is $O(2^n)$: the output is that large.

TSP is in a class of problems (NP-complete) for which **no polynomial algorithm is known**, and none may exist.

Exact is not always the goal

Is exhaustive TSP **correct**? Yes — if it finishes.

How much **time**? $O(n!)$ — it will not finish.

Can we **do better**? Better *exact* methods exist and are still exponential. We can also accept **approximate** tours.

A fast answer that is a bit long can beat a perfect answer that never arrives.


A greedy tour

1. Start at a city.
2. Go to the **closest unvisited** city.
3. Repeat until every city is visited; return home.

```
1 def nn_tour(dist, start):
2     n = len(dist)
3     visited = [start]
4     now = start
5     for _ in range(n - 1):
6         cand = [i for i in range(n) if i not in visited]
7         now = min(cand, key=lambda j: dist[now][j])
8         visited.append(now)
9     return visited
```

Nearest neighbor on the toy

Always the closest unused city. Same lab matrix as before. On reads, this was “glue the largest unused overlap.”



Start at A: $A \rightarrow B$ (2), $B \rightarrow D$ (1), $D \rightarrow C$ (4), home $C \rightarrow A$ (6). Tour $A-B-D-C-A$, length 13.

Start at B: $B \rightarrow D$ (1), $D \rightarrow A$ (3), $A \rightarrow C$ (6), home $C \rightarrow B$ (3). Tour $B-D-A-C-B$, length 13.

Same length, **different edges** (not a rotation of the same cycle). The shortest listed tour is still 12, so nearest neighbor missed the optimum either way.

Cost of nearest neighbor

At step k you scan the $n - k$ unused cities for the nearest.


$$T(n) = (n - 1) + (n - 2) + \cdots + 1 = O(n^2)$$

Method	Exact?	Cost
Exhaustive	yes	$O(n!)$
Nearest neighbor	no	$O(n^2)$

How good is it?

On random points in the plane, nearest neighbor is on average about **25% longer** than optimal (Johnson & McGeoch). That average hides ugly maps.


Six sampling sites around a long lake. Distance is a **straight line** (a crow, a boat). From Gunten, which unused site is closest?



Greedy sees only this step. Crossing is the closest hop, not a claim about the whole tour.

The leftover hop

Nearest neighbor takes that 5 km crossing, then walks on, and is left with Interlaken–Thun down the length of the lake.



The 24 km leftover is the bill for the 5 km crossing. Shore never takes that hop: 50 vs 59.


Greedy is $O(n^2)$ and useful. A locally nearest hop can be a globally stupid tour. Saving 5 km on one hop does not make a short tour — and the largest overlap need not assemble the genome.

Why uncross? 2-opt

Nearest neighbor is a **first guess**: a legal tour, built quickly. On the lab matrix it gave length **13**, but we already listed a tour of length **12**. So the guess can be improved without starting over.

Idea. Keep the same cities. Try a small change: drop two edges, reconnect their four endpoints the other way. If the new tour is shorter, keep the swap. That is **2-opt**.

On the toy, NN from *A* uses the two **diagonals**. Swap them for the two sides of the square: length 13 becomes 12.



Use NN (or any construction) as a starting tour, then search for improving swaps. Still no proof you have the shortest tour — only that this change helped.

Local search, not a guarantee

2-opt found the optimum on $n = 4$. That is luck of a tiny instance.


Local search is another *method for a good answer*. It still has no proof it is shortest, unless you check every tour.

In practice people run 2-opt (or 3-opt) after nearest neighbor. A genetic algorithm is another method, with an explicit representation: a tour is a permutation.

A genetic algorithm

A **population** of tours. Fitness = shortness (shorter is better). Make children; keep the short ones. Repeat for several generations.

1. **Select** two parents (prefer short tours).
2. **Crossover** that respects permutations — **order crossover**: copy a block from one parent; fill the rest in the other parent's order, skipping cities already used.
3. **Mutate** (swap two cities, or reverse a segment — a random 2-opt).



You choose the **parameters**: population size, number of generations, mutation rate. Different choices, different tours.

What the algorithm is — and is not

A genetic algorithm is a randomised search on a *population* of tours. It is not a model of evolution, and it has no optimality proof.

Ant colony, in one sentence: the same loop, but the object being mixed is **pheromone on edges**, not permutations of cities.

Simulated annealing: one tour; sometimes accept a worse swap so you can leave a local minimum.

Open the notebook: genetic algorithm

LAB BREAK · GENETIC ALGORITHM · PYTHON

Open `labs/02-tsp.ipynb` . Go to the **genetic algorithm** section (15 cities in the plane).

1. Run it. Press ► on the **animation** of the best tour (length curve beside it).
2. Change `POP_SIZE` , `GENERATIONS` , `MUTATION_RATE` , or `N_CITIES` . Run again.
3. Compare the final length to nearest neighbor (dashed line on the plot).

Play with the knobs. The lecture continues with a different greedy problem — one where the local rule **is** exact.

Methods that finish — a summary

Method	What it does	Exact?
Nearest neighbor	Always the closest unused city	No
2-opt	Uncross two edges if the tour shortens	No
Genetic algorithm	Population of tours; crossover + mutate	No
Ant colony / annealing	Same loop, different representation	No

We list $n = 4$ so you can see the method. None of these hops replaces listing tours if you need a proof.

Listing tours is exact and may never finish. A hop (NN, GA, 2-opt) can be wrong — and can still be the right tool when n is large.

Greedy can also be exact


TSP greedy can be wrong. That does not mean greedy is a bad idea. For some problems, the local rule **is** a proof.

Interval scheduling

The problem. One seminar room. Each talk has a start time and a finish time. Two talks cannot overlap; a talk **may start** when the previous one finishes. You want **as many talks as possible** — not the longest total time, the largest *count*.

A sequencing instrument that runs one library at a time is the same problem: start–finish windows, one machine.

One room. Talks have start and finish times. Overlaps cannot both be chosen.



Goal: as many talks as possible. A and C fit together; B overlaps both.

Input / output

INPUT

A list of talks. Each talk has a start time and a finish time.



OUTPUT


A largest set of talks with no two overlapping.

On the toy: three talks, times in hours from noon. A is $[0, 4]$, B is $[3, 5]$, C is $[4, 8]$.

Earliest finish

Always take the remaining talk that **finishes first**. Delete anything that overlaps it. Repeat.

Three talks, one room. A and C do not overlap. B is short and sits on the join.



Earliest finish: A then C (two talks).

Shortest first: B only. The local rule matters.

The local rule matters

Shortest talk first takes B (duration 2) and then nothing else overlaps. One talk.

Earliest finish takes A (finishes at 4), then C (starts at 4). Two talks.

Same idea (greedy), two rules. For this problem, earliest finish is **always** a largest set. Shortest first is not. For TSP, nearest neighbor is greedy and can be wrong.

Why earliest finish works

Exchange idea. Suppose some optimal schedule O disagrees with earliest-finish. Look at the first place they differ.

Earliest-finish picked talk A , which finishes as early as possible among remaining talks. Any optimal schedule must pick *some* talk that overlaps the early part of the day; that talk finishes **no earlier** than A .

Swap: replace that talk in O by A . The rest of O still fits, because A freed the room at least as soon. So there is an optimal schedule that agrees with earliest-finish on this choice. Repeat.

On the toy: earliest-finish takes A (done at 4), then C — two talks. If you take B instead, B occupies $[3, 5]$ and blocks both A and C — only one talk. Finishing early never blocks a larger set afterwards.

Open the notebook: interval scheduling

LAB · INTERVAL SCHEDULING · PYTHON

Still in `labs/02-tsp.ipynb`. Go to the **interval scheduling** section.

Both rules are already coded: `earliest_finish` and `shortest_first`. Pass any list of `(name, start, finish)` intervals.

1. Run both on the lecture toy A, B, C .
2. Run both on the **Q3 table** (hours 9–18) in the notebook.
3. Optional: invent your own list. Which rule keeps more talks?

Do not look for Q3 answers yet — they are at the end of the slide deck.

Exhaustive vs greedy

Exhaustive methods can return a correct optimum and may never finish. Greedy methods finish and may return a suboptimal answer — unless you have a proof for that rule (interval scheduling) or accept approximation (TSP).

Next week adds a third option: dynamic programming — still exact, much less work than listing everything, when subproblems overlap.

A bit of history

- **Menger** (1930s): the “messenger problem” — visit cities, return, minimise distance.
- **Dantzig, Fulkerson, Johnson** (1954): 49 US cities, cutting planes. Exact TSP becomes a research field.
- **Held & Karp** (1962): DP in $O(n^2 2^n)$. Still exponential; the one exact method worth naming.
- **Karp** (1972): Hamiltonian cycle is NP-complete. No known polynomial algorithm; none may exist.
- **Holland** (1975): genetic algorithms. A TSP tour is a permutation; order crossover is the later standard encoding.

Tools that use this

- Routing libraries (Google **OR-Tools**, **LKH**): tours and local search, not genomes.
- **Overlap-layout-consensus** assemblers (classical Celera, Canu for long reads): the layout step **visits every read once** — today's tour. Week 7 will draw this as a different graph.
- Circuit-board drilling and genome-probe ordering: the same tour, different labels.

Extensions

- **3-opt:** the same uncross, three edges at a time. The usual way NN is improved in practice after 2-opt.
- Week 5 will put exhaustive TSP back as a **search tree** with a bound.
- Interval scheduling is Kleinberg–Tardos, ch. 4 — the exchange argument we sketched.

Takeaway / Questions

Q1 How many distinct TSP tours on n cities, and why is that not $n!$?

Q2 Nearest neighbor from two starts can disagree, and neither need be shortest.

Q3 One room. As many talks as possible; two talks cannot overlap. A talk may start when the previous one finishes.

Talk	Start	Finish
<i>A</i>	09:00	13:00
<i>B</i>	12:00	14:00
<i>C</i>	13:00	17:00
<i>D</i>	17:00	18:00

- a. Listing every subset: how many subsets of 4 talks? Does that method give a largest possible set?
- b. Repeatedly take the remaining talk that **finishes first**. Which talks?
- c. Repeatedly take the remaining **shortest** talk that still fits. Which talks?
- d. Which of (a)–(c) is guaranteed a largest set? Which can return too few?

Leave Q3 open for now. In the lab notebook, both greedy rules are implemented — run them on this table and decide. Answers are at the **end** of these slides.

Takeaways

$n!$ grows faster than 2^n . Exhaustive TSP does not scale — neither does listing every order of n reads.

Nearest neighbor: always the closest unused city (always the largest unused overlap). $O(n^2)$, start-dependent, not always optimal. 2-opt can uncross; it still does not prove optimality.

Earliest-finish on intervals *is* optimal. Greedy is a method; the rule needs a proof or a counterexample.

Next week

Dynamic programming — the idea.

A tourist on a grid may only walk south and east. Listing every path is hopeless. The same cell is reached by many paths — so we will **fill a table** instead of recomputing.

To read

- Held & Karp (1962). A dynamic program for TSP, $O(n^2 2^n)$.
- Karp (1972). Reducibility among combinatorial problems — Hamiltonian cycle.
- Holland (1975): genetic algorithms; order crossover for permutations.
- Watch a tour grow: [TSP demo \(Google OR-Tools\)](#).

Q3 answers

After you have tried both rules in the lab.

Talk	Start	Finish	Duration
<i>A</i>	09:00	13:00	4 h
<i>B</i>	12:00	14:00	2 h
<i>C</i>	13:00	17:00	4 h
<i>D</i>	17:00	18:00	1 h

- (a) There are $2^4 = 16$ subsets. Checking every feasible subset and keeping the largest **is** exact — it finds a largest possible set. Cost grows as 2^n talks.
- (b) **Earliest finish:** sort by finish time *A*, *B*, *C*, *D*. Take *A* (ends 13:00), skip *B* (overlaps), take *C* (starts 13:00), take *D* (starts 17:00). Talks: *A*, *C*, *D* (three).
- (c) **Shortest first:** sort by duration *D* (1 h), *B* (2 h), then *A* or *C*. Take *D*, then *B* (no overlap with *D*). *A* and *C* both overlap *B*. Talks: *B*, *D* (only two).
- (d) Exact listing (a) and earliest finish (b) both give a largest set (size 3). Shortest first (c) can return too few (size 2 here). Same moral as the lecture toy: greedy is a method; **this** rule needs a